

Lab Report 3

Gradient Descent

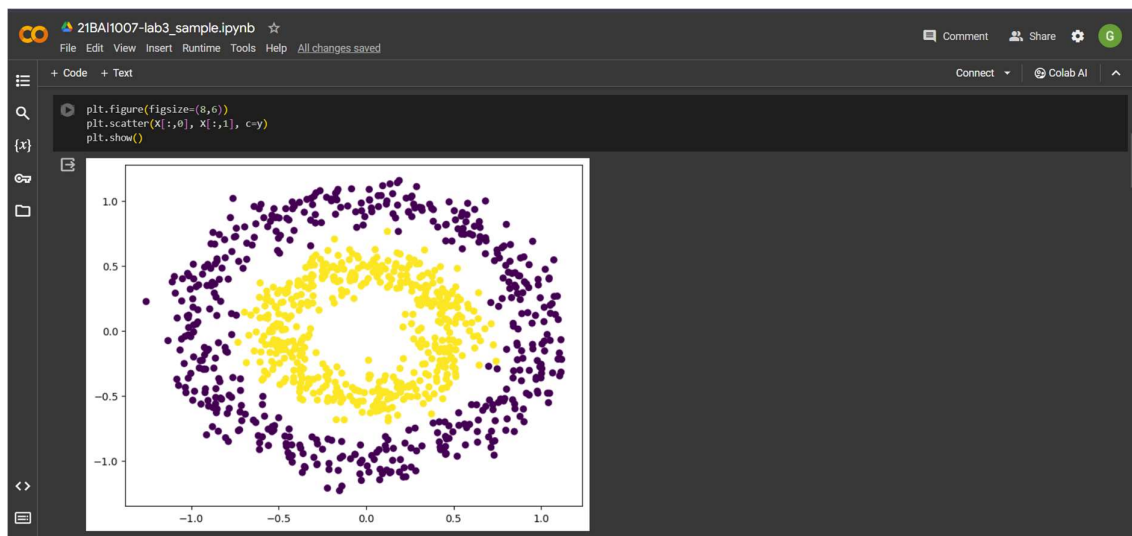
Goutham Krishnan

21BAI1007

Aim

1. Implement 3 different neural networks as listed below
 - 2 layers. 1st layer : ReLU (5 perceptrons), 2nd Layer : Sigmoid (1 perceptron)
 - In the above layers, both sigmoid
 - 4 layers with first three layers : Sigmoid (3) and the last layer : sigmoid (1)
2. Change the activation functions and implement the code given.
3. Implement a code to compare GD, SGD, Minibatch SGD for any input/neural model of your choice.

Q1. Implementing 3 different neural networks (sigmoid, ReLU)



```
21BAI1007-lab3_sample.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
1st Model - 1 ReLU(5), 1 Sigmoid(1) - 21BAI1007

[ ] model = Sequential()
model.add(Dense(5, activation='relu', input_dim=2))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

model.fit(X, y, epochs=100, verbose=1)

Epoch 1/100
32/32 [=====] - 0s 3ms/step - loss: 0.4422 - accuracy: 0.9480
Epoch 2/100
32/32 [=====] - 0s 3ms/step - loss: 0.4365 - accuracy: 0.9490
Epoch 3/100
32/32 [=====] - 0s 3ms/step - loss: 0.4309 - accuracy: 0.9490
Epoch 4/100
32/32 [=====] - 0s 3ms/step - loss: 0.4255 - accuracy: 0.9510
Epoch 5/100
32/32 [=====] - 0s 3ms/step - loss: 0.4199 - accuracy: 0.9560
Epoch 6/100
32/32 [=====] - 0s 3ms/step - loss: 0.4145 - accuracy: 0.9550
Epoch 7/100
32/32 [=====] - 0s 3ms/step - loss: 0.4089 - accuracy: 0.9560
Epoch 8/100
32/32 [=====] - 0s 3ms/step - loss: 0.4037 - accuracy: 0.9610
Epoch 9/100
32/32 [=====] - 0s 4ms/step - loss: 0.3983 - accuracy: 0.9610
Epoch 10/100
```

```
21BAI1007-lab3_sample.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[ ] print("Accuracy: ",model.evaluate(X, y)[1]*100,"%")
32/32 [=====] - 0s 1ms/step - loss: 0.1354 - accuracy: 0.9920
Accuracy: 99.19999837875366 %

2nd model - 1 sigmoid(5), 1 sigmoid(1) - 21BAI1007

[ ] model2 = Sequential()
model2.add(Dense(5, activation='sigmoid', input_dim=2))
model2.add(Dense(1, activation='sigmoid'))
model2.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

[ ] model2.fit(X, y, epochs=100, batch_size=32, verbose=1)

32/32 [=====] - 0s 2ms/step - loss: 0.6945 - accuracy: 0.5760
Epoch 10/100
32/32 [=====] - 0s 2ms/step - loss: 0.6944 - accuracy: 0.5370
Epoch 11/100
32/32 [=====] - 0s 2ms/step - loss: 0.6943 - accuracy: 0.5240
Epoch 12/100
32/32 [=====] - 0s 2ms/step - loss: 0.6943 - accuracy: 0.5190
Epoch 13/100
32/32 [=====] - 0s 2ms/step - loss: 0.6942 - accuracy: 0.5090
Epoch 14/100
32/32 [=====] - 0s 2ms/step - loss: 0.6941 - accuracy: 0.5170
Epoch 15/100
32/32 [=====] - 0s 2ms/step - loss: 0.6941 - accuracy: 0.5130
Epoch 16/100
32/32 [=====] - 0s 2ms/step - loss: 0.6940 - accuracy: 0.5120
Epoch 17/100
32/32 [=====] - 0s 2ms/step - loss: 0.6940 - accuracy: 0.5210
Epoch 18/100
```

```
21BAI1007-lab3_sample.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[ ] print("Accuracy: ",model2.evaluate(X, y)[1]*100,"%")
32/32 [=====] - 0s 1ms/step - loss: 0.6931 - accuracy: 0.5000
Accuracy: 50.0 %

3rd model - 3 sigmoid(3), 1 sigmoid(1) - 21BAI1007

[ ] model3 = Sequential()
model3.add(Dense(3, activation='sigmoid', input_dim=2))
model3.add(Dense(3, activation='sigmoid'))
model3.add(Dense(3, activation='sigmoid'))
model3.add(Dense(1, activation='sigmoid'))
model3.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

[ ] model3.fit(X, y, epochs=100, verbose=1)

Epoch 1/100
32/32 [=====] - 2s 4ms/step - loss: 0.6957 - accuracy: 0.5000
Epoch 2/100
32/32 [=====] - 0s 4ms/step - loss: 0.6945 - accuracy: 0.5000
Epoch 3/100
32/32 [=====] - 0s 4ms/step - loss: 0.6938 - accuracy: 0.5000
Epoch 4/100
32/32 [=====] - 0s 3ms/step - loss: 0.6934 - accuracy: 0.5000
Epoch 5/100
32/32 [=====] - 0s 3ms/step - loss: 0.6934 - accuracy: 0.5000
Epoch 6/100
32/32 [=====] - 0s 4ms/step - loss: 0.6935 - accuracy: 0.5000
Epoch 7/100
```

```
21BAI1007-lab3_sample.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

32/32 [=====] - 0s 2ms/step - loss: 0.6932 - accuracy: 0.4960
Epoch 24/100
32/32 [=====] - 0s 3ms/step - loss: 0.6932 - accuracy: 0.4970
Epoch 25/100
32/32 [=====] - 0s 2ms/step - loss: 0.6933 - accuracy: 0.4820
Epoch 26/100
32/32 [=====] - 0s 2ms/step - loss: 0.6932 - accuracy: 0.4730
Epoch 27/100
32/32 [=====] - 0s 2ms/step - loss: 0.6934 - accuracy: 0.4960
Epoch 28/100
32/32 [=====] - 0s 2ms/step - loss: 0.6931 - accuracy: 0.5000
Epoch 29/100
32/32 [=====] - 0s 2ms/step - loss: 0.6932 - accuracy: 0.4920

[ ] print("Accuracy: ",model3.evaluate(X, y)[1]*100,"%")

32/32 [=====] - 1s 1ms/step - loss: 0.6900 - accuracy: 0.6050
Accuracy: 60.50000190734863 %

[ ] Start coding or generate with AI.
```

Q1 extension. Implementing 3 different neural networks (sigmoid, ReLU) for the diabetes.csv Pima dataset

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

Implement 3 different neural networks as given below on the PIMA indian diabetes dataset -
21BAI1007

(i) 2 layers - ReLU(5), Sigmoid(1) (ii) 2 layers - Sigmoid(5), Sigmoid(1) (iii) 4 layers - Sigmoid(3), Sigmoid(3), Sigmoid(3), Sigmoid(1)

[ ] import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.callbacks import Callback
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras import Sequential
from tensorflow.keras.initializers import RandomNormal
import matplotlib.pyplot as plt
from sklearn.datasets import make_circles
from sklearn.metrics import accuracy_score

[ ] df = pd.read_csv('diabetes.csv')

[ ] len(df)

768
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

df.head(30)

Pregnancies Glucose BloodPressure SkinThickness Insulin BMI DiabetesPedigreeFunction Age Outcome
0 6 148 72 35 0 33.6 0.627 50 1
1 1 85 66 29 0 26.6 0.351 31 0
2 8 183 64 0 0 23.3 0.672 32 1
3 1 89 66 23 94 28.1 0.167 21 0
4 0 137 40 35 168 43.1 2.288 33 1
5 5 116 74 0 0 25.6 0.201 30 0
6 3 78 50 32 88 31.0 0.248 26 1
7 10 115 0 0 0 35.3 0.134 29 0
8 2 197 70 45 543 30.5 0.158 53 1
9 8 125 96 0 0 0.0 0.232 54 1
10 4 110 92 0 0 37.6 0.191 30 0
11 10 168 74 0 0 38.0 0.537 34 1
12 10 139 80 0 0 27.1 1.441 57 0
13 1 189 60 23 846 30.1 0.398 59 1
14 5 166 72 19 175 25.8 0.587 51 1
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

[ ] df.columns
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
      dtype='object')

[ ] df['Outcome'].value_counts()
0      500
1      268
Name: Outcome, dtype: int64

Processing the data and splitting the dataset into features(x) and labels(y)- 21BAI1007

[ ] x = df.iloc[:, :-1]
    y = df.iloc[:, -1:]

[ ] x.columns
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
      'BMI', 'DiabetesPedigreeFunction', 'Age'],
      dtype='object')

[ ] y.columns
Index(['Outcome'], dtype='object')
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

[ ] from sklearn.preprocessing import MinMaxScaler
    scaler = MinMaxScaler()
    x = scaler.fit_transform(np.array(x))

x
array([[0.35294118, 0.74371859, 0.59016393, ..., 0.50074516, 0.23441503,
        0.48333333],
       [0.05882353, 0.42713568, 0.54098361, ..., 0.39642325, 0.11656704,
        0.16666667],
       [0.47058824, 0.91959799, 0.52459016, ..., 0.34724292, 0.25362938,
        0.18333333],
       ...,
       [0.29411765, 0.6080402 , 0.59016393, ..., 0.390462 , 0.07130658,
        0.15 ],
       [0.05882353, 0.63316583, 0.49180328, ..., 0.4485842 , 0.11571307,
        0.43333333],
       [0.05882353, 0.46733668, 0.57377049, ..., 0.45305514, 0.10119556,
        0.63333333]])

[ ] x.shape
(768, 8)
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

1st Model - 1 ReLU(5), 1 Sigmoid(1) - 21BAI1007

[ ] model = Sequential()
    model.add(Dense(5, activation='relu', input_dim=x.shape[1]))
    model.add(Dense(1, activation='sigmoid'))
    model.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

model.fit(x, y, epochs=300, verbose=1)

Epoch 1/300
24/24 [=====] - 0s 3ms/step - loss: 0.5300 - accuracy: 0.7305
Epoch 2/300
24/24 [=====] - 0s 3ms/step - loss: 0.5271 - accuracy: 0.7253
Epoch 3/300
24/24 [=====] - 0s 3ms/step - loss: 0.5262 - accuracy: 0.7279
Epoch 4/300
24/24 [=====] - 0s 3ms/step - loss: 0.5251 - accuracy: 0.7305
Epoch 5/300
24/24 [=====] - 0s 3ms/step - loss: 0.5241 - accuracy: 0.7279
Epoch 6/300
24/24 [=====] - 0s 3ms/step - loss: 0.5227 - accuracy: 0.7292
Epoch 7/300
24/24 [=====] - 0s 3ms/step - loss: 0.5217 - accuracy: 0.7279
Epoch 8/300
24/24 [=====] - 0s 3ms/step - loss: 0.5207 - accuracy: 0.7266
Epoch 9/300
24/24 [=====] - 0s 2ms/step - loss: 0.5198 - accuracy: 0.7292
Epoch 10/300
24/24 [=====] - 0s 2ms/step - loss: 0.5186 - accuracy: 0.7292
Epoch 11/300
24/24 [=====] - 0s 2ms/step - loss: 0.5183 - accuracy: 0.7318
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

[ ] print("Accuracy: ",model.evaluate(x, y)[1]*100,"%")

24/24 [=====] - 0s 1ms/step - loss: 0.4412 - accuracy: 0.7969
Accuracy: 79.6875 %

2nd model - 1 sigmoid(5), 1 sigmoid(1) - 21BAI1007

[ ] model2 = Sequential()
model2.add(Dense(5, activation='sigmoid', input_dim=x.shape[1]))
model2.add(Dense(1, activation='sigmoid'))
model2.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

[ ] model2.fit(x, y, epochs=100, batch_size=32, verbose=1)

24/24 [=====] - 0s 2ms/step - loss: 0.6088 - accuracy: 0.6589
Epoch 73/100
24/24 [=====] - 0s 2ms/step - loss: 0.6000 - accuracy: 0.6589
Epoch 74/100
24/24 [=====] - 0s 2ms/step - loss: 0.5993 - accuracy: 0.6589
Epoch 75/100
24/24 [=====] - 0s 2ms/step - loss: 0.5985 - accuracy: 0.6576
Epoch 76/100
24/24 [=====] - 0s 2ms/step - loss: 0.5977 - accuracy: 0.6602
Epoch 77/100
24/24 [=====] - 0s 2ms/step - loss: 0.5970 - accuracy: 0.6602
Epoch 78/100
24/24 [=====] - 0s 2ms/step - loss: 0.5964 - accuracy: 0.6602
Epoch 79/100
24/24 [=====] - 0s 2ms/step - loss: 0.5955 - accuracy: 0.6576
Epoch 80/100
24/24 [=====] - 0s 2ms/step - loss: 0.5947 - accuracy: 0.6576
Epoch 81/100
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

[ ] Epoch 98/100
24/24 [=====] - 0s 2ms/step - loss: 0.5811 - accuracy: 0.6784
Epoch 99/100
24/24 [=====] - 0s 2ms/step - loss: 0.5804 - accuracy: 0.6758
Epoch 100/100
24/24 [=====] - 0s 2ms/step - loss: 0.5796 - accuracy: 0.6758
<keras.src.callbacks.History at 0x7ee410219f60>

[ ] print("Accuracy: ",model2.evaluate(x, y)[1]*100,"%")

24/24 [=====] - 0s 1ms/step - loss: 0.5791 - accuracy: 0.6771
Accuracy: 67.7083134651184 %

3rd model - 3 sigmoid(3), 1 sigmoid(1) - 21BAI1007

model3 = Sequential()
model3.add(Dense(3, activation='sigmoid', input_dim=x.shape[1]))
model3.add(Dense(3, activation='sigmoid'))
model3.add(Dense(3, activation='sigmoid'))
model3.add(Dense(1, activation='sigmoid'))
model3.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```

```
21BAI1007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
Connect Colab AI

model3.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')

[ ] model3.fit(x, y, epochs=100, verbose=1)

Epoch 24/100
24/24 [=====] - 0s 2ms/step - loss: 0.6465 - accuracy: 0.6510
Epoch 25/100
24/24 [=====] - 0s 2ms/step - loss: 0.6465 - accuracy: 0.6510
Epoch 26/100
24/24 [=====] - 0s 2ms/step - loss: 0.6464 - accuracy: 0.6510
Epoch 27/100
24/24 [=====] - 0s 2ms/step - loss: 0.6465 - accuracy: 0.6510
Epoch 28/100
24/24 [=====] - 0s 2ms/step - loss: 0.6468 - accuracy: 0.6510
Epoch 29/100
24/24 [=====] - 0s 2ms/step - loss: 0.6464 - accuracy: 0.6510
Epoch 30/100
24/24 [=====] - 0s 2ms/step - loss: 0.6464 - accuracy: 0.6510
Epoch 31/100
24/24 [=====] - 0s 2ms/step - loss: 0.6464 - accuracy: 0.6510
Epoch 32/100
24/24 [=====] - 0s 2ms/step - loss: 0.6463 - accuracy: 0.6510
Epoch 33/100
24/24 [=====] - 0s 2ms/step - loss: 0.6463 - accuracy: 0.6510
Epoch 34/100
24/24 [=====] - 0s 2ms/step - loss: 0.6463 - accuracy: 0.6510
Epoch 35/100
24/24 [=====] - 0s 2ms/step - loss: 0.6463 - accuracy: 0.6510
Epoch 36/100
24/24 [=====] - 0s 2ms/step - loss: 0.6462 - accuracy: 0.6510
```

```
21BA11007-diabetes-Dataset.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
Comment Share 6

+ Code + Text
Connect Colab AI

[ ] Epoch 41/100
24/24 [=====] - 0s 2ms/step - loss: 0.6461 - accuracy: 0.6510
Epoch 42/100
24/24 [=====] - 0s 2ms/step - loss: 0.6461 - accuracy: 0.6510
Epoch 43/100
24/24 [=====] - 0s 2ms/step - loss: 0.6461 - accuracy: 0.6510
Epoch 44/100
24/24 [=====] - 0s 2ms/step - loss: 0.6461 - accuracy: 0.6510
Epoch 45/100
24/24 [=====] - 0s 2ms/step - loss: 0.6461 - accuracy: 0.6510
Epoch 46/100
24/24 [=====] - 0s 2ms/step - loss: 0.6459 - accuracy: 0.6510
Epoch 47/100
24/24 [=====] - 0s 2ms/step - loss: 0.6460 - accuracy: 0.6510
Epoch 48/100
24/24 [=====] - 0s 2ms/step - loss: 0.6458 - accuracy: 0.6510
Epoch 49/100
24/24 [=====] - 0s 2ms/step - loss: 0.6458 - accuracy: 0.6510
Epoch 50/100
24/24 [=====] - 0s 2ms/step - loss: 0.6458 - accuracy: 0.6510
Epoch 51/100
24/24 [=====] - 0s 2ms/step - loss: 0.6457 - accuracy: 0.6510
Epoch 52/100
24/24 [=====] - 0s 2ms/step - loss: 0.6457 - accuracy: 0.6510

[ ] print("Accuracy: ",model3.evaluate(x, y)[1]*100,"%")

24/24 [=====] - 0s 1ms/step - loss: 0.6303 - accuracy: 0.6510
Accuracy: 65.10416865348816 %
```

Q2. Visualizing gradient descent for tanh and ReLU activation functions

```
21BA11007-GradientDescent.ipynb ☆
File Edit View Insert Runtime Tools Help Last edited on 4 February
Comment Share 6

+ Code + Text
Connect Colab AI

Visualizing gradient descent by building a model using tanh and ReLU activation functions

21BA11007

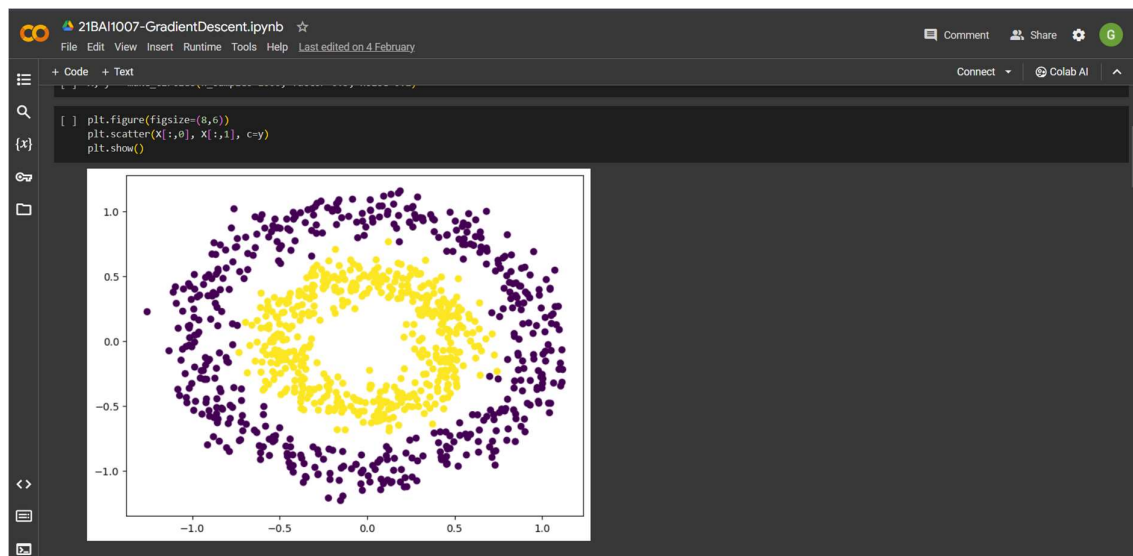
import numpy as np
import tensorflow as tf
from tensorflow.keras.callbacks import Callback
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.models import Sequential
from tensorflow.keras.initializers import RandomNormal
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score

[ ] from sklearn.datasets import make_circles

tf.random.set_seed(42)
np.random.seed(42)

X, y = make_circles(n_samples=1000, factor=0.5, noise=0.1)

[ ] plt.figure(figsize=(8,6))
plt.scatter(X[:,0], X[:,1], c=y)
plt.show()
```



```
21BA11007-GradientDescent.ipynb ☆
File Edit View Insert Runtime Tools Help Last edited on 4 February

+ Code + Text
Connect Colab AI

[] # illustrate weights across epochs
class WeightCapture(callback):
    "capture the weights of each layer of the model"
    def __init__(self, model):
        super().__init__()
        self.model = model
        self.weights = []
        self.epochs = []

    def on_epoch_end(self, epoch, logs=None):
        self.epochs.append(epoch) # remember the epoch axis
        weight = {}
        for layer in model.layers:
            if not layer.weights:
                continue
            name = layer.weights[0].name.split("/")[-1]
            weight[name] = layer.weights[0].numpy()
            self.weights.append(weight)

    def make_mlp(activation, initializer, name):
        "create a model with specified activation and initializer"
        model = Sequential([
            Input(shape=(2,)),
            Dense(5, activation=activation, kernel_initializer=initializer, name=name+"1"),
            Dense(5, activation=activation, kernel_initializer=initializer, name=name+"2"),
            Dense(5, activation=activation, kernel_initializer=initializer, name=name+"3"),
            Dense(5, activation=activation, kernel_initializer=initializer, name=name+"4"),
            Dense(1, activation="sigmoid", kernel_initializer=initializer, name=name+"5")
        ])
        return model
```

```
21BA11007-GradientDescent.ipynb ☆
File Edit View Insert Runtime Tools Help Last edited on 4 February

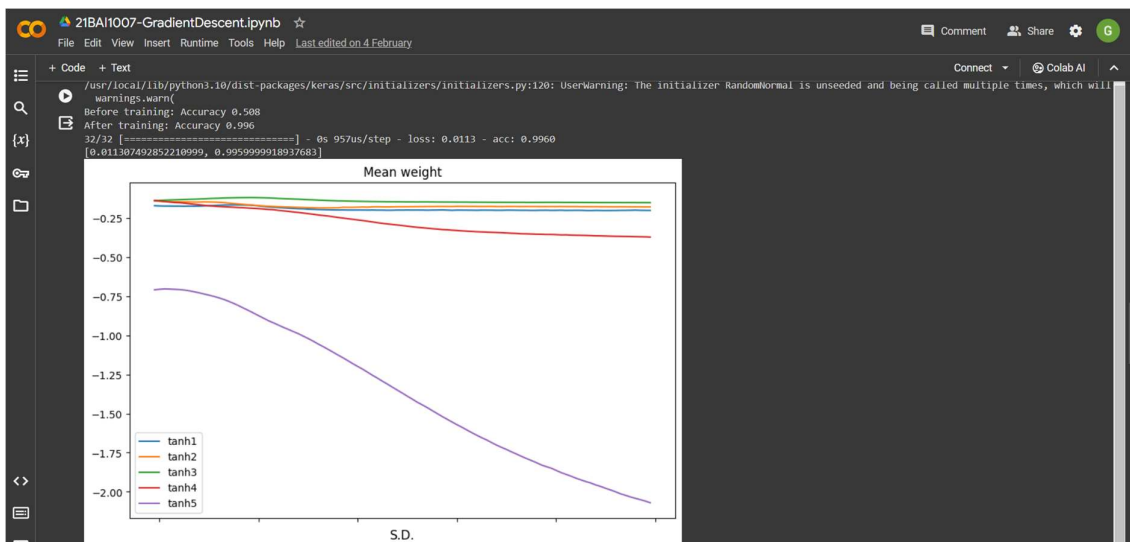
+ Code + Text
Connect Colab AI


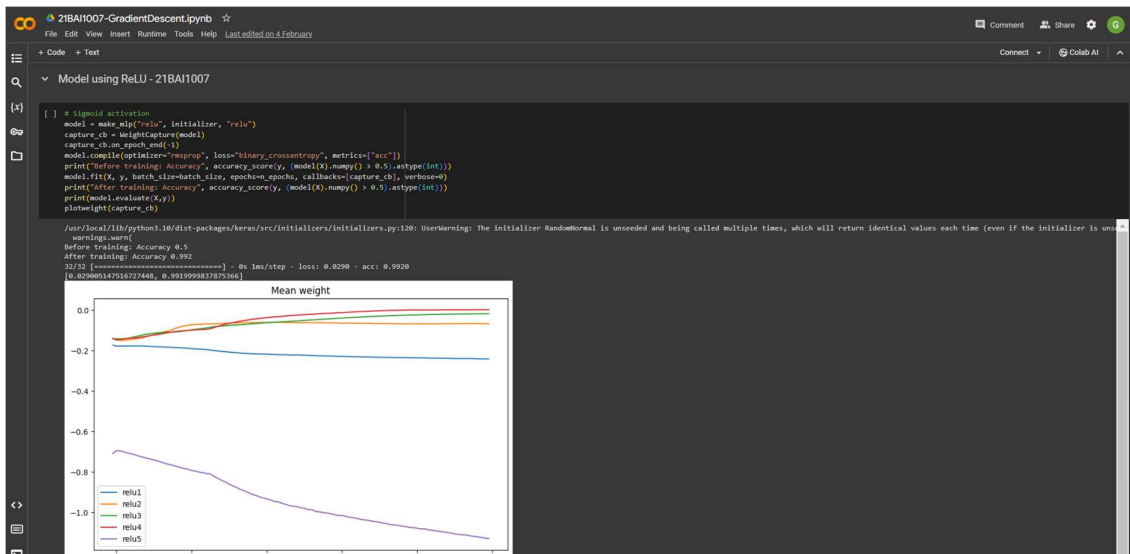
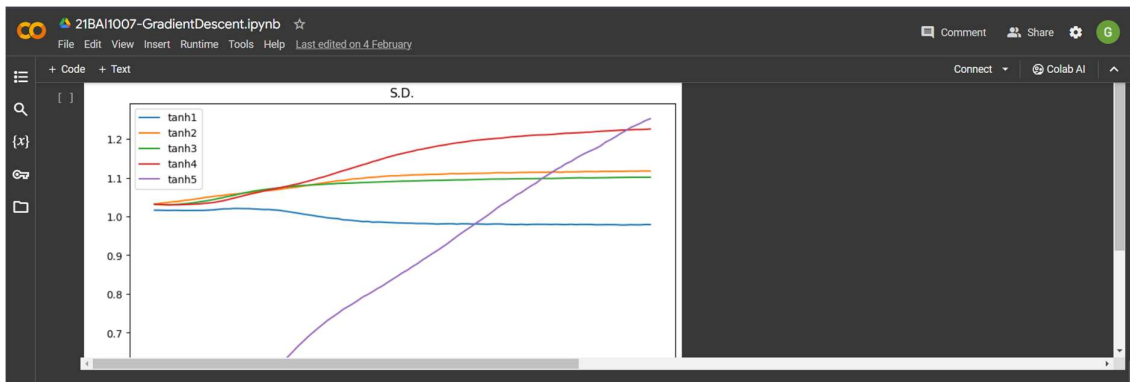
[] def plotweight(capture_cb):
    "Plot the weights' mean and s.d. across epochs"
    fig, ax = plt.subplots(2, 1, sharex=True, constrained_layout=True, figsize=(8, 10))
    ax[0].set_title("Mean weight")
    for key in capture_cb.weights[0]:
        ax[0].plot(capture_cb.epochs, [w[key].mean() for w in capture_cb.weights], label=key)
    ax[0].legend()
    ax[1].set_title("S.D.")
    for key in capture_cb.weights[0]:
        ax[1].plot(capture_cb.epochs, [w[key].std() for w in capture_cb.weights], label=key)
    ax[1].legend()
    plt.show()

[] initializer = RandomNormal(mean=0, stddev=1)
batch_size = 32
n_epochs = 100

Model using tanh - 21BA11007

[] # Sigmoid activation
model = make_mlp("tanh", initializer, "tanh")
capture_cb = WeightCapture(model)
capture_cb.on_epoch_end(-1)
model.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["acc"])
print("Before training: Accuracy", accuracy_score(y, (model(X).numpy() > 0.5).astype(int)))
model.fit(X, y, batch_size=batch_size, epochs=n_epochs, callbacks=[capture_cb], verbose=0)
print("After training: Accuracy", accuracy_score(y, (model(X).numpy() > 0.5).astype(int)))
print(model.evaluate(X,y))
plotweight(capture_cb)
```





Q3. Implementing Gradient Descent, Stochastic gradient descent and mini batch SGD

```
21BAI1007-q3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
RAM Disk Colab AI
Implementing GD, SGD and Mini Batch SGD
21BAI1007 Goutham Krishnan
[25] import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras.callbacks import Callback
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.models import Sequential
from tensorflow.keras.initializers import RandomNormal
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

[26] df = pd.read_csv('diabetes.csv')

[27] df.shape
(768, 9)

[28] df.head()


|   | Pregnancies | Glucose | BloodPressure | Skinthickness | Insulin | BMI  | DiabetesPedigreefunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6           | 148     | 72            | 35            | 0       | 33.6 | 0.627                    | 50  | 1       |
| 1 | 1           | 85      | 66            | 29            | 0       | 26.6 | 0.351                    | 31  | 0       |
| 2 | 8           | 183     | 64            | 0             | 0       | 23.3 | 0.672                    | 32  | 1       |
| 3 | 1           | 89      | 66            | 23            | 94      | 28.1 | 0.167                    | 21  | 0       |
| 4 | 0           | 137     | 40            | 35            | 168     | 43.1 | 2.288                    | 33  | 1       |


```

```
21BAI1007-q3.ipynb
File Edit View Insert Runtime Tools Help All changes saved
RAM Disk Colab AI
Cleaning and processing diabetes dataset - 21BAI1007
[29] x = df.iloc[:,1:-1]
y = df.iloc[:, -1:]

[30] X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=42)

Creating and training the model - 21BAI1007
[31] from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(np.array(X_train))
X_test_scaled = scaler.transform(np.array(X_test))

[32] X_train_scaled_bias = np.c_[np.ones((len(X_train_scaled), 1)), X_train_scaled]
X_test_scaled_bias = np.c_[np.ones((len(X_test_scaled), 1)), X_test_scaled]

[33] def classification_model(input_dim):
model = tf.keras.models.Sequential([
tf.keras.layers.Input(shape=(input_dim,)),
tf.keras.layers.Dense(16, activation='relu'),
tf.keras.layers.Dense(8, activation='relu'),
tf.keras.layers.Dense(1, activation='sigmoid')
])
return model

[34] def train_model(optimizer, X_train, y_train, epochs=100, batch_size=None):
model = classification_model(input_dim=X_train.shape[1])
model.compile(optimizer=optimizer, loss='binary_crossentropy')
history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size, verbose=0)
return model, history.history['loss']

Results from the model - 21BAI1007
[55] gd_model, gd_loss = train_model(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01), X_train=X_train_scaled_bias, y_train=y_train)

[56] sgd_model, sgd_loss = train_model(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01), X_train=X_train_scaled_bias, y_train=y_train, batch_size=1)

[57] minibatch_sgd_model, minibatch_sgd_loss = train_model(optimizer=tf.keras.optimizers.SGD(learning_rate=0.01), X_train=X_train_scaled_bias, y_train=y_train, batch_size=32)

[58] gd_test_loss = gd_model.evaluate(X_test_scaled_bias, y_test, verbose=0)
sgd_test_loss = sgd_model.evaluate(X_test_scaled_bias, y_test, verbose=0)
minibatch_sgd_test_loss = minibatch_sgd_model.evaluate(X_test_scaled_bias, y_test, verbose=0)

[59] print("GD Test loss:", gd_test_loss)
print("SGD Test loss:", sgd_test_loss)
print("Mini-batch SGD Test loss:", minibatch_sgd_test_loss)

GD Test loss: 0.5320150256160921
SGD Test loss: 0.5012460076608818
Mini-batch SGD Test loss: 0.6190852522850017

[60] import matplotlib.pyplot as plt

# Plot the training loss curves
plt.figure(figsize=(10, 6))
plt.plot(gd_loss, label='GD')
plt.plot(sgd_loss, label='SGD')
plt.plot(minibatch_sgd_loss, label='Mini-batch SGD')
plt.title('Training Loss Curves')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.show()
```

