

Lab Report – 7

Autoencoders, Object Detection

Goutham Krishnan

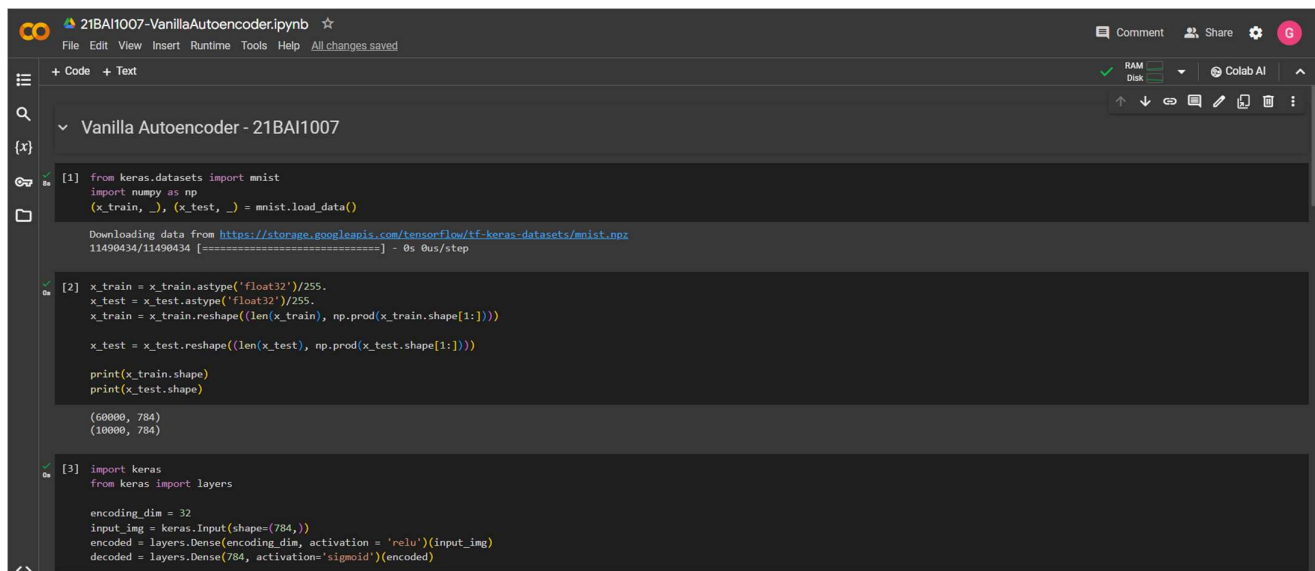
21BAI1007

Aim

1. Execute the two sample programs, one for vanilla autoencoder and one for denoising autoencoder on the MNIST dataset.
2. Create an object detection model using YOLO that detects the eyes and face.
3. Using autoencoders, implement the dimensionality reduction of MNIST handwritten image dataset.

Observations and Output

For the sample code



```
21BAI1007-VanillaAutoencoder.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
RAM
Disk
Colab AI
6

Vanilla Autoencoder - 21BAI1007

[1] from keras.datasets import mnist
import numpy as np
(x_train, _), (x_test, _) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

[2] x_train = x_train.astype('float32')/255.
x_test = x_test.astype('float32')/255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

print(x_train.shape)
print(x_test.shape)

(60000, 784)
(10000, 784)

[3] import keras
from keras import layers

encoding_dim = 32
input_img = keras.Input(shape=(784,))
encoded = layers.Dense(encoding_dim, activation='relu')(input_img)
decoded = layers.Dense(784, activation='sigmoid')(encoded)
```

```
21BA11007-VanillaAutoencoder.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
autoencoder = keras.Model(input_img, decoded)
encoder=keras.Model(input_img, encoded)

encoded_input = keras.Input(shape=(encoding_dim,))

decoder_layer = autoencoder.layers[-1]

decoder = keras.Model(encoded_input, decoder_layer(encoded_input))
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.fit(x_train, x_train, epochs=50, batch_size=256, shuffle=True, validation_data=(x_test, x_test))


235/235 [=====] - 3s 12ms/step - loss: 0.0933 - val_loss: 0.0920
Epoch 23/50
235/235 [=====] - 3s 12ms/step - loss: 0.0932 - val_loss: 0.0920
Epoch 24/50
235/235 [=====] - 3s 11ms/step - loss: 0.0931 - val_loss: 0.0919
Epoch 25/50
235/235 [=====] - 2s 10ms/step - loss: 0.0931 - val_loss: 0.0919
Epoch 26/50
235/235 [=====] - 3s 11ms/step - loss: 0.0930 - val_loss: 0.0919
Epoch 27/50
235/235 [=====] - 3s 13ms/step - loss: 0.0930 - val_loss: 0.0919
Epoch 28/50
235/235 [=====] - 3s 11ms/step - loss: 0.0930 - val_loss: 0.0918
Epoch 29/50
235/235 [=====] - 2s 11ms/step - loss: 0.0929 - val_loss: 0.0917
Epoch 30/50
235/235 [=====] - 2s 11ms/step - loss: 0.0929 - val_loss: 0.0917
Epoch 31/50
235/235 [=====] - 2s 11ms/step - loss: 0.0929 - val_loss: 0.0916
Epoch 32/50
235/235 [=====] - 3s 14ms/step - loss: 0.0928 - val_loss: 0.0917
Epoch 33/50
235/235 [=====] - 2s 10ms/step - loss: 0.0928 - val_loss: 0.0917
Epoch 34/50
235/235 [=====] - 2s 10ms/step - loss: 0.0928 - val_loss: 0.0917
Epoch 35/50
235/235 [=====] - 2s 10ms/step - loss: 0.0928 - val_loss: 0.0916
```


```
21BA11007-VanillaAutoencoder.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
encoded_imgs = encoder.predict(x_test)
decoded_imgs = decoder.predict(encoded_imgs)

313/313 [=====] - 1s 2ms/step
313/313 [=====] - 0s 1ms/step

[6] import matplotlib.pyplot as plt
n=10
plt.figure(figsize=(20, 4))

for i in range(n):
    ax = plt.subplot(2, n, i+1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    plt.imshow(decoded_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



```
21BA11007-VanillaAutoencoder.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[6] 
```

```
21BA1007-Denoising_Autoencoder.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Denoising Autoencoders - 21BA1007

from keras.datasets import mnist
import numpy as np
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.astype('float32')/255.
x_test = x_test.astype('float32')/255.

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

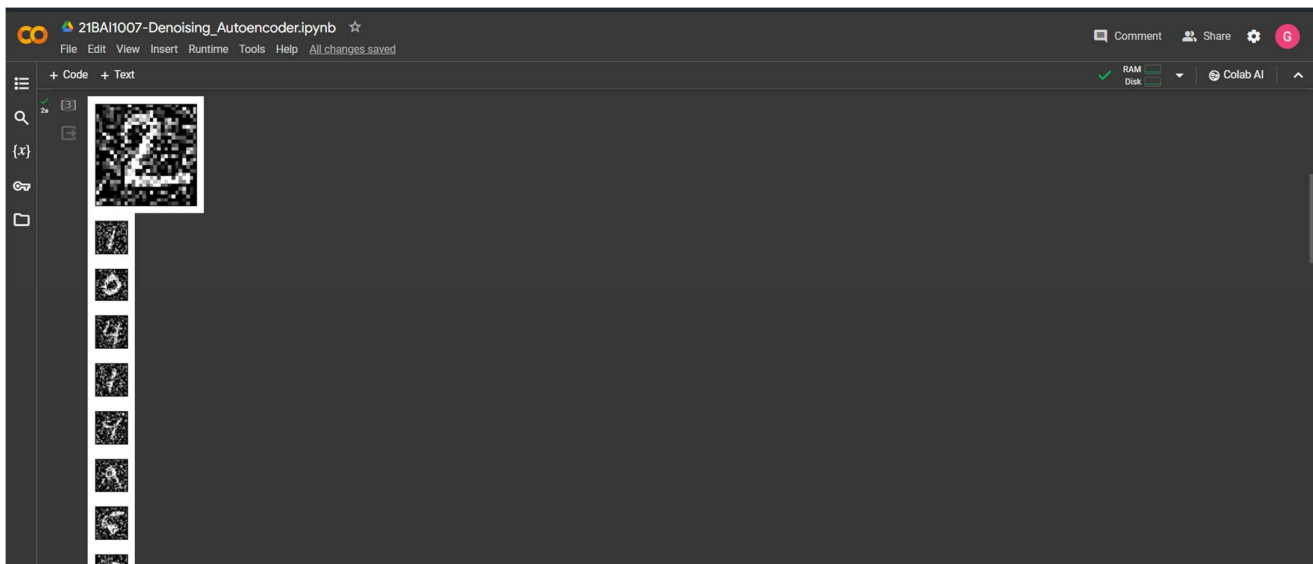
[2] x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))

noise_factor = 0.5

x_train_noisy = x_train + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_train.shape)
x_test_noisy = x_test + noise_factor * np.random.normal(loc=0.0, scale=1.0, size=x_test.shape)
x_train_noisy = np.clip(x_train_noisy, 0., 1.)
x_test_noisy = np.clip(x_test_noisy, 0., 1.)

[3] import matplotlib.pyplot as plt
n = 10
plt.figure(figsize=(20, 2))

for i in range(1, n+1):
    ax = plt.subplot(1, n, i)
    plt.imshow(x_test_noisy[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()
```



```
21BA1007-Denoising_Autoencoder.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text

[4] import keras
from keras import layers
import tensorflow as tf
from keras.callbacks import TensorBoard
encoding_dim = 32
input_img = keras.Input(shape=(784,))

encoded = layers.Dense(encoding_dim, activation='relu')(input_img)
decoded = layers.Dense(784, activation='sigmoid')(encoded)

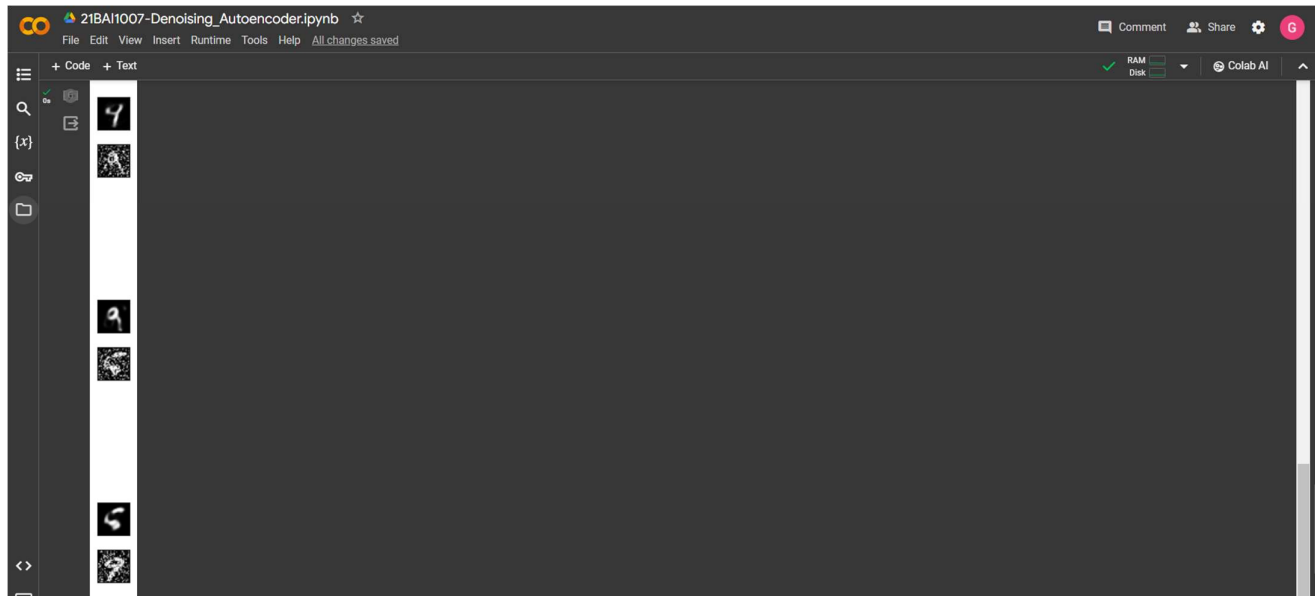
autoencoder = keras.Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.fit(x_train_noisy, x_train, epochs=100, batch_size=128, shuffle=True, validation_data=(x_test_noisy, x_test),)
encoder = keras.Model(input_img, encoded)

Epoch 72/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1247
Epoch 73/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1247
Epoch 74/100
469/469 [=====] - 4s 9ms/step - loss: 0.1255 - val_loss: 0.1246
Epoch 75/100
469/469 [=====] - 3s 7ms/step - loss: 0.1255 - val_loss: 0.1246
Epoch 76/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1245
Epoch 77/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1245
Epoch 78/100
469/469 [=====] - 4s 9ms/step - loss: 0.1255 - val_loss: 0.1245
Epoch 79/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1249
Epoch 80/100
469/469 [=====] - 3s 6ms/step - loss: 0.1255 - val_loss: 0.1246
Epoch 81/100
469/469 [=====] - 3s 7ms/step - loss: 0.1255 - val_loss: 0.1247
Epoch 82/100
469/469 [=====] - 4s 9ms/step - loss: 0.1254 - val_loss: 0.1245
Epoch 83/100
```

```
21BA1007-Denoising_Autoencoder.ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[5] encoded_input = keras.Input(shape = (encoding_dim,))
    decoder_layer = autoencoder.layers[-1]
    decoder = keras.Model(encoded_input, decoder_layer(encoded_input))
    encoded_imgs = encoder.predict(x_test_noisy)
    decoded_imgs = decoder.predict(encoded_imgs)

313/313 [=====] - 1s 2ms/step
313/313 [=====] - 1s 2ms/step

[6] import matplotlib.pyplot as plt
    n = 10
    plt.figure(figsize=(20, 4))
    for i in range(n):
        ax = plt.subplot(2, n, i+1)
        plt.imshow(x_test_noisy[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
        ax = plt.subplot(2, n, i+1+n)
        plt.imshow(decoded_imgs[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)
    plt.show()
```



For Face Detection Model using YOLO

Code:

```
from ultralytics import YOLO
import cv2
import math
# start webcam
cap = cv2.VideoCapture(0)
cap.set(3, 640)
cap.set(4, 480)

# model
model = YOLO("yolov8n-face.pt")

# object classes
classNames = ["face"]

while True:
    success, img = cap.read()
    results = model(img, stream=True)
```

```

# coordinates
for r in results:
    boxes = r.boxes

    for box in boxes:
        # bounding box
        x1, y1, x2, y2 = box.xyxy[0]
        x1, y1, x2, y2 = int(x1), int(y1), int(x2), int(y2) # convert
to int values

        # put box in cam
        cv2.rectangle(img, (x1, y1), (x2, y2), (255, 0, 255), 3)

        # confidence
        confidence = math.ceil((box.conf[0]*100))/100
        print("Confidence --->", confidence)

        # class name
        cls = int(box.cls[0])
        print("Class name -->", classNames[cls])

        # object details
        org = [x1, y1]
        font = cv2.FONT_HERSHEY_SIMPLEX
        fontScale = 1
        color = (255, 0, 0)
        thickness = 2


        cv2.putText(img, classNames[cls], org, font, fontScale, color,
thickness)

    cv2.imshow('Webcam', img)
    if cv2.waitKey(1) == ord('q'):
        break

cap.release()
cv2.destroyAllWindows()

```

Output:



For dimensionality reduction using Autoencoders

```
21BA11007-DimReduction-MNIST.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
T4 RAM Disk Colab AI

Dimensionality Reduction of MNIST Dataset using AutoEncoders
Goutham Krishnan 21BA11007

[1] import numpy as np
import matplotlib.pyplot as plt
from keras.datasets import mnist
from keras.models import Model, Sequential
from keras.layers import Reshape, Flatten, Dense, Lambda
from keras import losses

Loading the data

[2] (train_images, _), (test_images, _) = mnist.load_data()

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 0s 0us/step

Scaling the data

[3] # Scaling
x_train = train_images.astype('float32')/255.
x_test = test_images.astype('float32')/255.

print(x_train.shape)
print(x_test.shape)

(60000, 28, 28)
(10000, 28, 28)
```

```
21BA11007-DimReduction-MNIST.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text
T4 RAM Disk Colab AI

Creating the Autoencoder model

latent_dim = 64

class Autoencoder(Model):
    def __init__(self, latent_dim):
        super(Autoencoder, self).__init__()
        self.latent_dim = latent_dim
        self.encoder = Sequential([
            Flatten(),
            Dense(latent_dim, activation='relu'),
        ])
        self.decoder = Sequential([
            Dense(784, activation='sigmoid'),
            Reshape((28, 28)),
        ])

    def call(self, x):
        encoded = self.encoder(x)
        decoded = self.decoder(encoded)
        return decoded

autoencoder = Autoencoder(latent_dim)

[5] autoencoder.compile(optimizer='adam', loss=losses.MeanSquaredError())
```

```
21BA11007-DimReduction-MNIST.ipynb ☆
File Edit View Insert Runtime Tools Help Saving...

+ Code + Text
T4 RAM Disk Colab AI

Train the model

autoencoder.fit(x_train, x_train, epochs=20, shuffle=True)

Epoch 1/20
1875/1875 [=====] - 8s 3ms/step - loss: 0.0239
Epoch 2/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0068
Epoch 3/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0050
Epoch 4/20
1875/1875 [=====] - 5s 2ms/step - loss: 0.0045
Epoch 5/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0043
Epoch 6/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0042
Epoch 7/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0042
Epoch 8/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0041
Epoch 9/20
1875/1875 [=====] - 5s 2ms/step - loss: 0.0040
Epoch 10/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0040
Epoch 11/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0040
Epoch 12/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0039
Epoch 13/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0039
Epoch 14/20
1875/1875 [=====] - 5s 3ms/step - loss: 0.0039
Epoch 15/20
1875/1875 [=====] - 4s 2ms/step - loss: 0.0039
Epoch 16/20
1875/1875 [=====] - 5s 2ms/step - loss: 0.0039
```

```
21BA1007-DimReduction-MNIST.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
[7] print(autoencoder.encoder.summary())

Model: "sequential"
Layer (type) Output Shape Param #
-----
flatten (Flatten) (32, 784) 0
dense (Dense) (32, 64) 50240
-----
Total params: 50240 (196.25 KB)
Trainable params: 50240 (196.25 KB)
Non-trainable params: 0 (0.00 Byte)
None

print(autoencoder.decoder.summary())

Model: "sequential_1"
Layer (type) Output Shape Param #
-----
dense_1 (Dense) (32, 784) 50960
reshape (Reshape) (32, 28, 28) 0
-----
Total params: 50960 (199.06 KB)
Trainable params: 50960 (199.06 KB)
Non-trainable params: 0 (0.00 Byte)
None
```

```
21BA1007-DimReduction-MNIST.ipynb ☆
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
Reconstruct the test images
encoded_imgs = autoencoder.encoder(test_images).numpy()
decoded_imgs = autoencoder.decoder(encoded_imgs).numpy()

[10] n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    ax = plt.subplot(2, n, i+1)
    plt.imshow(test_images[i])
    plt.title("Original images")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    ax = plt.subplot(2, n, i+1+n)
    plt.imshow(decoded_imgs[i])
    plt.title("Reconstructed images")
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

plt.show()
```

